

Chapter 1: Introduction

Objectives:

- To provide a grand tour of the major operating systems components
- To provide coverage of basic computer system organization

What Operating Systems Do

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating System goals
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
- Use the computer hardware in an efficient manner
- Four Components of a Computer System
 - Computer hardware
 - CPU
 - Memory
 - I/O devices
 - Operating system
 - Controls and coordinates use of hardware among various applications and users
 - system and application programs
 - Define the ways in which the system resources are used to solve the computing problems of the users
 - word processors, compilers, web browsers, database systems, video games
 - Users
 - People, machines, other computers
- OS Definition
- resource allocator
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- control program
 - Controls execution of programs to prevent errors and improper use of the computer

Computer-System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles
- Common Functions of Interrupts
 - Interrupt transfers control to the interrupt service routine generally, through the interrupt vector, which contains the addresses of all the service routines
 - Interrupt architecture must save the address of the interrupted instruction

- Incoming interrupts are disabled while another interrupt is being processed to prevent a lost interrupt
 - A **trap** is a software-generated interrupt caused either by an error or a user request
 - An operating system is interrupt driven
- Interrupt Handling
- I/O structure
- I/O methods
 - Synchronous
 - Asynchronous
- Device-Status Table
- Direct Memory Access Structure
- Storage Structure
 - registers
 - cache
 - Main memory
 - electronic disk
 - magnetic disk
 - optical disk
 - magnetic tapes
- Why Storage is organized in hierarchy
 - Speed
 - Cost
 - Volatility
- Caching

Computer-System Architecture

Operating-System Structure

- Multiprogramming needed for efficiency
 - Multi-user
 - organizes jobs
 - a subset of total jobs in system is kept in memory
 - job scheduling
 - job switching
- Timesharing (multitasking) is logic extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating interactive computing.
 - response time
 - process
 - CPU scheduling
 - swapping
 - virtual memory

Operating-System Operations

- Interrupt driven by hardware
- software error or request creates exception or trap
- other process problems including infinite loop, ...
- Dual-mode operation allows OS to protect itself and other system components
 - User mode and kernel mode
 - Mode bit provided by hardware
 - bit to distinguish when system is running user code or kernel code
 - designated as privileged
 - system call can change mode to kernel and return from call resets it to users.

Process Management

- A process is a program in execution. It is a unit of work within the system. **Program** is a passive entity, **process** is an active entity.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one program counter specifying location of next instruction to execute
- Multi-threaded process has one program counter per thread
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
- The operating system is responsible for the following activities in connection with process management:
 - Creating and deleting both user and system processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling

Memory Management

- All data in memory before and after processing
- All instructions in memory in order to execute
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit, **file**
 - Each medium is controlled by device (i.e., disk drive, tape drive)
- File-System management

Protection and Security

- Protection – any mechanism for controlling access of processes or users to resources defined by the OS
- Security – defense of the system against internal and external attacks
- Systems generally first distinguish among users, to determine who can do what

Computing Environments

- Traditional computer
 - Office environment
 - Home networks
- Client-Server Computing
 - Dumb terminals supplanted by smart PCs
 - Many systems now servers, responding to requests generated by clients
- Peer-to-Peer Computing
 - Another model of distributed system
 - P2P does not distinguish clients and servers
- Web-Based Computing

Chapter 2: Operating-System Structure

Objectives:

- To describe the services an operating system provides to users, processes, and other systems
- To discuss the various ways of structuring an operating system
- To explain how operating systems are installed and customized and how they boot

Operating System Services

One set of operating-system services provides functions that are helpful to the user:

- User interface
- Program execution
- I/O operations
- File-system manipulation
- Communications
- Error detection
- Another set of OS functions exists for ensuring the efficient operation of the system itself via resource sharing
 - Resource allocation
 - Accounting

- Protection and security

System Calls

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level Application Program Interface (API) rather than direct system call use

Types of System Calls

- Process control
- File management
- Device management
- Information maintenance
- Communications

System Programs

System programs provide a convenient environment for program development and execution. They can be divided into:

- File manipulation
- Status information
- File modification
- Programming language support
- Program loading and execution
- Communications
- Application programs

Layered Approach

Layered Operating System

Microkernel System Structure

- Moves as much from the kernel into “user” space
- Communication takes place between user modules using message passing
- Benefits:
 - Easier to extend a microkernel
 - Easier to port the operating system to new architectures
 - More reliable (less code is running in kernel mode)
 - More secure
- Detriments:
 - Performance overhead of user space to kernel space communication

Modules

- Most modern operating systems implement kernel modules
- Overall, similar to layers but with more flexible

Operating System Generation

- Operating systems are designed to run on any of a class of machines
- SYSGEN program
- Booting
- Bootstrap program

Chapter 3: Processes

Process Concept

- An operating system executes a variety of programs:
 - Batch system
 - Time-shared systems
- Process – a program in execution; process execution must progress in sequential fashion
- A process includes:
 - program counter
 - stack
 - data section
- As a process executes, it changes state
 - new
 - running
 - waiting
 - ready
 - terminated

Process Control Block (PCB)

- Information associated with each process
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information

Process Scheduling Queues

- Job queue
- Ready queue
- Device queues
- Processes migrate among the various queues

Schedulers

- Long-term scheduler (or job scheduler)
- Short-term scheduler (or CPU scheduler)
- The long-term scheduler controls the degree of multiprogramming
- Processes can be described as either:

- I/O-bound process
- CPU-bound process

Context Switch

Process Creation

Process Termination

- Process executes last statement and asks the operating system to delete it (exit)
- Parent may terminate execution of children processes (abort)

Cooperating Processes

- Independent process cannot affect or be affected by the execution of another process
- Cooperating process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Producer-Consumer Problem

- unbounded-buffer places no practical limit on the size of the buffer
- bounded-buffer assumes that there is a fixed buffer size

Interprocess Communication (IPC)

- Direct Communication
- Indirect Communication

Synchronization

- Blocking is considered synchronous
- Non-blocking is considered asynchronous

Client-Server Communication

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

Chapter 4: Threads

Benefits

- Responsiveness
- Resource Sharing
- Economy
- Utilization of MP Architectures

User Threads

Kernel Threads

Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many

Two-level Model

- Similar to M:M, except that it allows a user thread to be bound to kernel thread

Threading Issues

- Semantics of fork() and exec() system calls
- Thread cancellation
 - Terminating a thread before it has finished
 - Two general approaches:
 - Asynchronous cancellation
 - Deferred cancellation
- Signal handling
- Thread pools
- Thread specific data

Chapter 5: CPU Scheduling

CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
 - Switches from running to waiting state
 - Switches from running to ready state
 - Switches from waiting to ready
 - Terminates
- Scheduling under 1 and 4 is non-preemptive
- All other scheduling is preemptive

Dispatcher

Scheduling Criteria

- CPU utilization
- Throughput
- Turnaround Time
- Waiting Time
- Average waiting time

- Response Time

First-Come, First-Served (FCFS) Scheduling

Shortest-Job-First (SJF) Scheduling

- Non-Preemptive SJF
- Preemptive SJF

Priority Scheduling

Round Robin (RR)

- Time Quantum

Multilevel Queue Scheduling

Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Multiple-Processor Scheduling

Real-Time Scheduling

- Hard real-time systems – required to complete a critical task within a guaranteed amount of time
- Soft real-time computing – requires that critical processes receive priority over less fortunate ones

Algorithm Evaluation

- Deterministic modeling
- Queueing models
- Implementation

Chapter 6: Process Synchronization

Solution to Critical-Section Problem

- Mutual Exclusion
- Progress Bounded Waiting

Peterson's Solution

Synchronization Hardware

TestAndSet Instruction

Swap Instruction

Semaphore

- Binary semaphore
- Counting semaphore

Deadlock and Starvation

Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Monitors

Chapter 7: Deadlocks

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set

Deadlock Characterization

- Mutual exclusion
- Hold and wait
- No preemption
- Circular wait

Resource Allocation Graph

- If graph contains no cycles
 - no deadlock
- If graph contains a cycle
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock

Methods for Handling Deadlocks

- Ensure that the system will never enter a deadlock state
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system

Deadlock Prevention

- Mutual Exclusion – not required for sharable resources; must hold for non-sharable resources.
- Hold and Wait – must guarantee that whenever a process requests a resource, it does not hold any other resources.
 - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none.
 - Low resource utilization; starvation possible.
- No Preemption
- Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.

Deadlock Avoidance

- Safe State
- If a system is in safe state \rightarrow no deadlocks.
- If a system is in unsafe state \rightarrow possibility of deadlock.
- Avoidance \rightarrow ensure that a system will never enter an unsafe state.

Avoidance algorithms

- Single instance of a resource type. Use a resource-allocation graph
- Multiple instances of a resource type. Use the banker's algorithm

Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
 - the same as Avoidance algorithms
- Recovery scheme

Recovery from Deadlock: Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort
 - Priority of the process
 - How long process has computed, and how much longer to completion
 - Resources the process has used
 - Resources process needs to complete
 - How many processes will need to be terminated
 - Is process interactive or batch

Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost
- Rollback – return to some safe state, restart process for that state
- Starvation – same process may always be picked as victim, include number of rollback in cost factor

Chapter 8: Memory Management

Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - Compile time
 - Load time
 - Execution time

Memory-Management Unit (MMU)

- Hardware device that maps virtual to physical address

Dynamic Loading

Dynamic Linking

Swapping

Contiguous Allocation

Dynamic Storage-Allocation Problem

- How to satisfy a request of size n from a list of free holes
 - First-fit
 - Best-fit
 - Worst-fit

Fragmentation

- External Fragmentation
- Internal Fragmentation
- Reduce external fragmentation by compaction

Paging

- Divide physical memory into fixed-sized blocks called **frames** (size is power of 2)
- Divide logical memory into blocks of same size called **pages**

Effective Access Time

- Associative Lookup = ϵ time unit
- Assume memory cycle time is 1 microsecond
- Hit ratio α – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} &= (1 + \epsilon)\alpha + (2 + \epsilon)(1 - \alpha) \\ &= 2 + \epsilon + \alpha \end{aligned}$$

Memory Protection

Structure of the Page Table

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

Chapter 9: Virtual Memory

- Virtual memory – separation of user logical memory from physical memory.
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation

Demand Paging

- Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Lazy swapper

Valid-Invalid Bit

Page Fault

Performance of Demand Paging

- Page Fault Rate $p \leq 1.0$
- Effective Access Time (EAT)

$$EAT = (1 - p)\alpha + p\beta$$

Copy-on-Write

- Copy-on-Write (COW) allows both parent and child processes to initially share the same pages in memory
- If either process modifies a shared page, only then is the page copied

Page Replacement

- First-In-First-Out (FIFO) Algorithm
 - Belady's Anomaly
- Optimal Algorithm
- Least Recently Used (LRU) Algorithm
 - Counter implementation
 - Stack implementation
- LRU Approximation Algorithms
 - Reference bit
 - Second chance

- Counting Algorithms
 - LFU Algorithm
 - MFU Algorithm

Allocation of Frames

- Fixed Allocation
- Priority Allocation
- Global vs. Local Allocation
 - Global replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another
 - Local replacement – each process selects from only its own set of allocated frames

Thrashing

- Why does thrashing occur
 - size of locality > total memory size

Working-Set Model

- Δ - working-set window - a fixed number of page references
- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
- $D = \sum WSS_i$ - total demand frames
- if $D > m \rightarrow$ Thrashing
- Policy if $D > m$, then suspend one of the processes

Other Issues -- Pre-paging

Other Issues -- Page Size

Other Issues – Program Structure

Other Issues – I/O interlock

Chapter 10: File-System Interface

- Contiguous logical address space

Open File Locking

File-system Organization

- Single-Level Directory
- Two-Level Directory
- Tree-Structured Directories
- Acyclic-Graph Directories

File System Mounting

File Sharing

Protection

Chapter 11: File System Implementation

Directory Implementation

- Linear list
- Hash Table

Allocation Methods

- Contiguous allocation
- Linked allocation
- Indexed allocation

Free-Space Management

- Bit vector
- Linked list
- Indexed list

Recovery

Chapter 12: Mass-Storage Systems

- Magnetic disks
 - provide bulk of secondary storage of modern computers
 - Disks can be removable
 - Drive attached to computer via I/O bus
- Magnetic tape

Disk Scheduling

- FCFS
- SSTF
- SCAN
- C-SCAN
- C-LOOK

Disk Management

- Low-level formatting, or physical formatting — Dividing a disk into sectors that the disk controller can read and write

Chapter 13: I/O Systems

I/O Hardware

- Incredible variety of I/O devices
- Common concepts
 - Port
 - Bus (daisy chain or shared direct access)
 - Controller (host adapter)
- I/O instructions control devices
- Devices have addresses, used by
 - Direct I/O instructions
 - Memory-mapped I/O

Polling

- Determines state of device
 - command-ready
 - busy
 - error
- Busy-wait cycle to wait for I/O from device

Interrupts

- CPU Interrupt-request line triggered by I/O device
- Interrupt handler receives interrupts
- Maskable to ignore or delay some interrupts
- Interrupt vector to dispatch interrupt to correct handler
 - Based on priority
 - Some non-maskable

Direct Memory Access

- Used to avoid programmed I/O for large data movement
- Requires DMA controller
- Bypasses CPU to transfer data directly between I/O device and memory

Application I/O Interface

A Kernel I/O Structure

Block and Character Devices

- Block devices include disk drives
- Character devices include keyboards, mice, serial ports

Network Devices

- Varying enough from block and character to have own interface
- Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)

Clocks and Timers

- Programmable interval timer used for timings, periodic interrupts

Blocking and Nonblocking I/O

- Blocking - process suspended until I/O completed
 - Easy to use and understand
 - Insufficient for some needs
- Nonblocking - I/O call returns as much as available
 - User interface, data copy (buffered I/O)
 - Implemented via multi-threading
 - Returns quickly with count of bytes read or written
- Asynchronous - process runs while I/O executes
 - Difficult to use
 - I/O subsystem signals process when I/O completed

Two I/O Methods

- Synchronous
- Asynchronous

Kernel I/O Subsystem

- Scheduling
- Buffering - store data in memory while transferring between devices
- Caching - fast memory holding copy of data
- Spooling - hold output for a device
- Device reservation - provides exclusive access to a device

Device-status Table

Error Handling

- OS can recover from disk read, device unavailable, transient write failures
- Most return an error number or code when I/O request fails
- System error logs hold problem reports

I/O Protection

- All I/O instructions defined to be privileged
- I/O must be performed via system calls

Kernel Data Structures

I/O Requests to Hardware Operations

Life Cycle of An I/O Request

Streams

Intercomputer Communications

Device-Functionality Progression